



gokul-1998 / MLP_sep_2023



<> Code

Issues

Pull requests

Actions

Projects

Wiki

Security

Insights

Settings

MLP_sep_2023 / oppe1 prep / mock.md



Gokulakrishnan more

6cdc868 · 17 hours ago



History



MOCK FOR OPPE

1) **Online Remote Proctored Exams**

1 point

This exam is being conducted in the online mode which will be taken by the examinee at his/her place of residence and proctored remotely by those assigned by the IIT Madras team. During such an exam, the exam session will be recorded by the proctors and the recordings will be handed over to IIT Madras which will be kept for internal purposes. The following guidelines are to be followed by all examinees.

- No examinee shall share their personal details with the proctors (including but not limited to their phone numbers, address), during or after the exam.
- The table/desk where the examinee takes this exam shall not have any items kept that may have sensitive information including but not limited to their phone numbers and address.
- No examinee shall aid, or attempt to aid another candidate by discussing answers via email, text etc.
- No examinee will disclose any of the details of what happened during the exam or examination trials to anyone outside.
- If an examinee wishes to ask a question during the exam, they should post the query in the exam room chat window and the proctor will clarify the issue.
- If any examinee is found to have violated any of the Code of Conduct for Online Examinations, or in any other way to have

[MLP_sep_2023](#) / [oppe1 prep](#) / [mock.md](#)[↑ Top](#)**Preview**

Code Blame 392 lines (285 loc) · 11.5 KB

Raw



Instructions**2) Instructions - MLP****1 point**

Login 30 minutes before the exam and complete your identity verification and room verification process

The latest version of seeya has to be downloaded and run. If seeya is not working a local recording has to be taken (Inform the proctor of this)

You will be provided with a Colab file for working on the dataset. The Colab file should have the relevant codes related to each question that you are attempting. We will be evaluating the colab file too for releasing the final marks.

You will not be allowed to log out of the exam meeting before the end of 120 minutes.

For each question, copy its description/body from the portal to the colab in its proper place.

Only scikit-learn, numpy, scipy, pandas and inbuilt python libraries are allowed to be used.

☐ Yes, I Agree

Preamble: Download breast cancer dataset from sklearn and answer the questions:1-5

3) How many features are there in the data set?

2 points

4) How many unique classes are there in the dataset?

2 points

5) How many samples correspond to class 0?

4 points

3.

```
from sklearn.datasets import load_breast_cancer

cancer = load_breast_cancer()
# How many features are there in the data set?
print(len(cancer.feature_names))
```



- Ans : 30

4.

```
from sklearn.datasets import load_breast_cancer

cancer = load_breast_cancer()
# How many unique classes are there in the dataset?
print((len(cancer.target_names)))
```



- Ans: 2

5.

```
from sklearn.datasets import load_breast_cancer

cancer = load_breast_cancer()
# How many samples correspond to class 0?
x=cancer.target[cancer.target == 0].shape[0]
print(x)
```



Ans : 212

6.

```
from sklearn.datasets import load_breast_cancer
```



```
X,y = load_breast_cancer(as_frame=True,return_X_y=True)
# Is there any categorical feature?
# we need to have it as pandas dataframe to use info()
print(X.info())
```

- Ans : No

7.

```
from sklearn.datasets import load_breast_cancer
```



```
X,y = load_breast_cancer(as_frame=True,return_X_y=True)

# What is the pearson correlation between the features 'mean perimeter' and 'mean radius'?
# Round your answer to 3 decimal places.
# Hint: Use the pandas.DataFrame.corr() method.

print(X.corr().loc['mean perimeter','mean radius'].round(3))
```

- Ans : 0.998

8.

```
# Import SimpleImputer from SkLearn library.
# # Take a matrix [[7, 16, 31], [np.nan, np.nan, 66], [12, 5, np.nan],[98, np.nan, 92]].
# # Impute the missing values in the matrix using SimpleImputer with Mean What is the imputed matrix
using fit_transform?
```



```
import numpy as np
from sklearn.impute import SimpleImputer

matrix = np.array([[7, 16, 31], [np.nan, np.nan, 66], [12, 5, np.nan], [98, np.nan, 92]])
# print(matrix)

imp = SimpleImputer(missing_values=np.nan, strategy='mean')
imp.fit(matrix)
matrix = imp.transform(matrix)
print(matrix)
```

- Ans : [[7. 16. 31.] [39. 10.5 66.] [12. 5. 63.] [98. 10.5 92.]]

9.



```
# Import FunctionTransformer from the SkLearn library.
# Apply log base 10 to the elements of the following array : [[1, 1], [2, 3], [10, 100]] .
# What is the resultant matrix?
```

```
from sklearn.preprocessing import FunctionTransformer
import numpy as np
```

```
arr = np.array([[1, 1], [2, 3], [10, 100]])
# print(arr)
```

```
transformer = FunctionTransformer(np.log10)
print(transformer.transform(arr))
```

- Ans : [[0. 0.] [0.30103 0.47712125] [1. 2.]]

10.



```
# Read the CSV file from https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data, define the column headers as `sepal length', 'sepal width', 'petal length', 'petal width','label'.
```

```
# Generate a new feature matrix consisting of all polynomial combinations of the features with degree 2 (For example, if an input sample is two dimensional and of the form [a,b] , the degree-2 polynomial features are
```

```
#
```

```
What are the shapes of the feature matrix before and after the polynomial transformation?
```

```
# What are the shapes of the feature matrix before and after the polynomial transformation?
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.preprocessing import PolynomialFeatures
```

```
df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data', header=None)
```

```
df.columns = ['sepal length', 'sepal width', 'petal length', 'petal width','label']
```

```
print(df.shape)
```

```
X = df[['sepal length', 'sepal width', 'petal length', 'petal width']]
```

```
y = df['label']
```

```
poly = PolynomialFeatures(2)
```

```
X_poly = poly.fit_transform(X)
```

```
print(X_poly.shape)
```

- Ans : (150, 5) (150, 15)

11.

```
# Import the California Housing dataset and SelectPercentile, mutual_info_regression.
```

```
# Select features according to 10 percentile of the highest scores What is the shape of the feature matrix after feature selection?
```



```
from sklearn.datasets import fetch_california_housing
X,y=fetch_california_housing(return_X_y=True,as_frame=True)
from sklearn.feature_selection import SelectPercentile, mutual_info_regression
sp=SelectPercentile(mutual_info_regression,percentile=10)
new=sp.fit_transform(X,y)
new.shape
```

- Ans : (20640, 1)

12.

```
# Import the California Housing dataset.
# Load the features and labels as numpy array. Split the data into training and test data in 4:1
proportion. What will be the size of training features?
```



```
from sklearn.datasets import fetch_california_housing
import numpy as np
from sklearn.model_selection import train_test_split

cal_housing = fetch_california_housing()
X = cal_housing.data
y = cal_housing.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

print(X_train.shape)
```

- Ans : (16512, 8)

13.



```
# Import the California Housing dataset.
# Load the features and labels as numpy array. Split the data into training and test data in 4:1
proportion. What will be the size of training features?

from sklearn.datasets import fetch_california_housing
import numpy as np
from sklearn.model_selection import train_test_split

cal_housing = fetch_california_housing()
X = cal_housing.data
y = cal_housing.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

print(X_train.shape)

# Use the strategy 'mean' to predict all the labels. Calculate the coefficient of determination (R2) for
the prediction.

from sklearn.dummy import DummyRegressor
from sklearn.metrics import r2_score

dummy = DummyRegressor(strategy='mean')
dummy.fit(X_train, y_train)
y_pred = dummy.predict(X_test)
print(r2_score(y_test, y_pred))
```

- Answer differs for every run, so find out why and how to get the desired answer

14.

```
# Import the California Housing dataset.
# Load the features and labels as numpy array. Split the data into training and test data in 4:1
proportion. What will be the size of training features?
```



```
from sklearn.datasets import fetch_california_housing
import numpy as np
from sklearn.model_selection import train_test_split

cal_housing = fetch_california_housing()
X = cal_housing.data
y = cal_housing.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

print(X_train.shape)

# Fit the training dataset using Iterative Optimization, i.e., SGDRegressor (Keep all the parameters
default) after scaling the features using MinMaxScaler.\

# Calculate the mean absolute error for training data.

from sklearn.linear_model import SGDRegressor
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_absolute_error

scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

sgd = SGDRegressor()
sgd.fit(X_train_scaled, y_train)

y_pred = sgd.predict(X_train_scaled)
print(mean_absolute_error(y_train, y_pred))
```

- Ans : 0.5601

15.



```
# Import the California Housing dataset.
# Load the features and labels as numpy array. Split the data into training and test data in 4:1
proportion. What will be the size of training features?

from sklearn.datasets import fetch_california_housing
import numpy as np
from sklearn.model_selection import train_test_split

cal_housing = fetch_california_housing()
X = cal_housing.data
y = cal_housing.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

print(X_train.shape)

# Fit the training dataset using Iterative Optimization, i.e., SGDRegressor (Keep all the parameters
default) after scaling the features using MinMaxScaler.\

# Calculate the mean absolute error for training data.

from sklearn.linear_model import SGDRegressor
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_absolute_error

scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

sgd = SGDRegressor()
sgd.fit(X_train_scaled, y_train)

y_pred = sgd.predict(X_train_scaled)
```

```
print(mean_absolute_error(y_train, y_pred))  
# Calculate the mean absolute error for test data.  
  
y_pred = sgd.predict(X_test_scaled)  
print(mean_absolute_error(y_test, y_pred))
```

- Ans : 0.5601 0.5601

16.

```
# Import the California Housing dataset.  
# Load the features and labels as numpy array. Split the data into training and test data in 4:1  
proportion. What will be the size of training features?  
  
from sklearn.datasets import fetch_california_housing  
import numpy as np  
from sklearn.model_selection import train_test_split  
  
cal_housing = fetch_california_housing()  
X = cal_housing.data  
y = cal_housing.target  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)  
  
print(X_train.shape)  
  
# Fit the training dataset using Iterative Optimization, i.e., SGDRegressor (Keep all the parameters  
default) after scaling the features using MinMaxScaler.  
  
# Calculate the mean absolute error for training data.  
  
from sklearn.linear_model import SGDRegressor  
from sklearn.preprocessing import MinMaxScaler
```



```
from sklearn.metrics import mean_absolute_error

scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

sgd = SGDRegressor()
sgd.fit(X_train_scaled, y_train)

y_pred = sgd.predict(X_train_scaled)
# print(mean_absolute_error(y_train, y_pred))
# Calculate the mean absolute error for test data.

y_pred = sgd.predict(X_test_scaled)
# print(mean_absolute_error(y_test, y_pred))

# print the scores that are calculated using 5 fold cross validation using LinearRegression on entire
dataset.

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score

lr = LinearRegression()
scores = cross_val_score(lr, X, y, cv=5)
print(scores)
```

- array([0.54866323, 0.46820691, 0.55078434, 0.53698703, 0.66051406])

17.

```
# Import the diabetes dataset from sklearn and check how many samples are there in the dataset?
```



```
from sklearn import datasets
```

```
import numpy as np
import pandas as pd

diabetes = datasets.load_diabetes()
print(diabetes.data.shape)
```

- Ans : (442, 10)

18.



```
from sklearn.datasets import load_diabetes

diabetes = load_diabetes()
# Step-1: Use Standard scaler to preprocess the data.
# Step-2: Split the dataset in such a way that 20% data is taken for test cases.(set random state=10)
# Step-3: Use the LinearRegression() estimator to predict the output.
# What is the R 2 score you got using LinearRegression estimator on test data.

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

scaler = StandardScaler()
X = scaler.fit_transform(diabetes.data)
y = diabetes.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=10)
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
print(r2_score(y_test, y_pred))
```

- Ans : 0.5341988244945842

19.



```
from sklearn.datasets import load_diabetes

diabetes = load_diabetes()
# Step-1: Use Standard scaler to preprocess the data.
# Step-2: Split the dataset in such a way that 20% data is taken for test cases.(set random state=10)
# Step-3: Use the LinearRegression() estimator to predict the output.
# What is the R 2 score you got using LinearRegression estimator on test data.
# What is the value of coefficient associated with variable "s3"?
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

scaler = StandardScaler()
X = scaler.fit_transform(diabetes.data)
y = diabetes.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=10)
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
print(lr.coef_[diabetes.feature_names.index('s3')])
```